



CUDA matmul size sweep


Runtime vs size



Throughput vs size



Result matrix



Collected metrics:

size	kernel_ms	gflops	duration_us	mem_pct	dram_pct	sm_pct	occ_pct	grid_size	block_size_reported
64	0.020070	26.122449	8.288000	10.565049	NaN	41.292404	16.472514	16.0	256.0
128	0.028314	148.137432	33.664000	20.410400	NaN	80.695943	42.463019	64.0	256.0
256	0.186571	179.847855	230.656000	23.580143	NaN	93.771133	82.848485	256.0	256.0
512	1.436357	186.886336	1.777696	24.367804	NaN	97.186488	95.250212	1024.0	256.0
1024	11.374029	188.805886	14.074240	24.563574	NaN	98.108754	98.407627	4096.0	256.0

Fastest runtime: 64x64x64 with 0.020 ms.

Best throughput: 1024x1024x1024 with 188.81 GFLOP/s.